

Cookies and Sessions

Lecture Notes for CS 142

Fall 2010

John Ousterhout

- Readings for this topic: none.

Stateless applications

- Web application servers are generally "stateless":
 - Each HTTP request is independent; server can't tell if 2 requests came from the same browser or user.
 - Web server applications maintain no information in memory from request to request (only information on disk survives from one request to another).
- Statelessness not always convenient for application developers: need to tie together a series of requests from the same user.

Browser cookies

- Cookie basics:
 - The first time a browser connects with a particular server, there are no cookies.
 - When the server responds it includes a `Set-Cookie:` header that defines a cookie.
 - Each cookie is just a name-value pair.
 - In the future whenever the browser connects with the same server, it includes a `Cookie:` header containing the name and value, which the server can use to connect related requests.
- What's in a cookie?
 - Name and data.
 - Data size limited by browsers (typically < 4 KB).
 - A server can define multiple cookies with different names, but browsers limit the number of cookies per server (around 50).
 - Domain for this cookie: server, port (optional), URL prefix (optional). The cookie is only included in requests matching its domain.
 - Expiration date: browser can delete old cookies.

Sessions

- Cookies are used by the server to implement *sessions*:
 - A pool of data related to an active connection (one browser instance).
- Typically the cookie for an application contains an identifier for a session.
- Web frameworks like Rails do most of the work of managing sessions and cookies:
 - Rails provides `session`, a hash-like object in which you can store anything you like
 - Data will be available in all future requests from the same browser.
 - Rails automatically checks for a session cookie at the start of each request:
 - Cookie exists? use it to find session data

- No cookie? Create new session, new cookie
 - End of each request: save session data where it can be found by future requests.
- Managing session state:
 - Approach #1: just keep state in main memory
 - Approach #2: store session state in files on disk
 - Approach #3: store session state in a database
 - Most frameworks allow you to control session storage:
 - Provide an object that saves and restores session data.
- Server must eventually delete stale session data.
- Sessions have numerous security issues, which we will discuss later.